

SUYASH

Data Scientist | AI Architect

[LinkedIn](#) | [GitHub](#) | [Portfolio](#)

suyashkumarthakur@gmail.com ❖ +1(469) 428-6430 ❖ Dallas, Texas, USA

Data Scientist & AI Architect with 8+ years delivering \$250M+ in validated business impact. I am either building 0-to-1 data science solutions or aggressively optimizing existing enterprise systems from 1 to 100. Known for finding novel solutions and owning the entire lifecycle —from uncovering the hidden insight to engineering the end-to-end production pipeline. Currently building AI Experimentation platform at Albertsons solo.

WORK EXPERIENCE

Albertsons

Feb 2024 – Present

Data Scientist IV(Lead Platform Developer)

Pleasanton, CA

- **Architected End-to-End Causal Engine:** Replaced manual workflows with a proprietary experimentation platform, scaling velocity from 3 to 8+ experiments/week and driving \$150M+ in annual validated decisions.
- **Engineered Advanced Statistical Pipeline:** Integrated Spark pipelines with CUPED variance reduction (35% error reduction) and Bayesian inference, compressing test cycles from 28 to 12 days.
- **Deployed GenAI Analyst Agent:** Designed a SQL+RAG system using LangChain and Gemini that translates natural language queries into verified SQL; reduced post-experiment analysis time from days to hours.
- **Operationalized Context Engineering:** Developed error-logging and context protocols for LLM agents, reducing hallucinations by 60% and saving stakeholders 8+ hours/week.
- **Automated Results Summarization:** Built a Beta Auto-Summary Agent that autonomously reads experiment outputs and generates stakeholder reports, eliminating the final manual bottleneck in the testing lifecycle.

Applied AI Research Lab

Jan 2023 – Present

Principal Engineer & Product Builder

Remote, US

- **Established Technical Sandbox:** Founded an applied research lab to validate emerging Agentic and RAG architectures before enterprise deployment.
- **QuizBeef (700+ Users):** Architected an AI-native quiz platform using LLM semantic parsing and FAISS-powered vector retrieval to generate quizzes from multiple documents; enabled multi-source question mixing, adaptive difficulty, and skill-based question selection inspired by spaced-repetition models. Served as a toy-lab for RAG and agentic patterns later applied to Albertsons' GenAI Copilot.
- **mana-health:** Developed a multi-agent health platform orchestrating specialized agents (physio, nutritionist) to process biometric and workout data.
- **ClinicOS:** AI-powered clinic management SaaS live in 8 clinics across rural India; handles real-time consultation transcription, auto-generated prescriptions, appointment queuing, and inventory management.

Discover Financial Services

Aug 2022 – Oct 2023

Senior Data Science Analyst

Chicago, IL

- **Overtaken decade-old risk policy:** Used XGBoost + Propensity Score Matching to prove employment verification had zero predictive value for high-FICO borrowers ($R^2 < 0.01$), then designed surgical A/B test that unlocked \$114M annualized volume and 300+ incremental customers/month through friction reduction.
- **Enabled \$1.8M in loan approvals :** Identified and resolved critical iOS bugs through funnel analysis on the Automated Loan Approval platform, collaborating with engineering and UI/UX teams to restore conversion flow.
- **Infrastructure Modernization:** Migrated legacy SAS workflows to GCP BigQuery, reducing infrastructure costs by 12% and improving decision speed by 10+ hours/week.

Vodafone

Jun 2016 – Feb 2019

Assistant Manager, Data Science

Pune, IN

- Increased marketing revenue by \$11.2M and improved CTR by 28% by orchestrating email personalization using logistic regression to optimize content targeting, then A/B testing new campaign strategies.
- Enhanced customer lifetime value by 5% by developing Linear Regression models to optimize subscription pricing based on user data and predicted retention patterns across customer segments.

Senior Executive, Data Science

- Enabled real-time performance tracking across 10+ models by partnering with product teams to define and instrument 15+ custom metrics, building SQL pipelines and executive dashboards for new feature evaluation.
- Reduced Priority 2 production issues by 60%+ by developing Tableau dashboards with time series forecasting alerts, enabling proactive incident detection and faster resolution.

Kaizen

Jun 2021 – Dec 2021

Data Science Engineering Intern

Atlanta, GA

- Informed 3+ marketing campaigns by analyzing user clustering patterns to uncover distinct customer journey types, translating data insights into actionable segmentation strategies.
- Reduced authentication false positives by 13% by building ETL pipelines and deploying an anomaly detection model on AWS SageMaker, adding risk-based flagging to enhance 2FA security.
- Improved forecasting accuracy by 8.5% by designing a hybrid time-series model integrating Prophet and LSTM for Toyota sales prediction, outperforming baseline statistical models.

EDUCATION

The University of Texas at Dallas

Jun 2022

MS, Business Analytics

Dallas, TX

- Dean's Scholar (Top 20 %); President, Data Science Club — led workshops and speaker series.
- Top Student Mentor for Data Science track subjects. Answered 800+ questions and coached 70+ students.

Pune University

Jun 2016

BS, Computer Science and Engineering

Pune, IN

SKILLS

GenAI & Agentic Systems: LangChain, RAG (Retrieval-Augmented Generation), Context Engineering, LLMs, Prompt Engineering, Autonomous Agents, SQL+RAG, dbt, Snowflake

Causal Inference & Experimentation: Causal Inference, CUPED, Bayesian Methods (PyMC), Propensity Score Matching, Synthetic Controls, Heterogeneous Treatment Effects (HTE), Uplift Modeling, A/B Testing, SRM Detection

Data & Infrastructure: Python, SQL, PySpark, BigQuery, Airflow, Spark, GCP (Vertex AI), AWS (S3, SageMaker), Docker, Pandas, NumPy, Scikit-learn, TensorFlow

Visualization & Tools: Tableau, Streamlit, Git, Azure

Certifications

[Google Analytics](#)